

Normalized Read Latency vs PAUSE count [normalized to master]
spinlock: [trylock→PAUSE×N]×30→mutex_lock / mc=4t mut -T4 -c1 -d32 / median across runs / ★ = QPS peak N per arch: Broadwell N=45, Ice N=15, Emerald N=30, Skylake N=5

